

# P9120 - Statistical Learning and Data Mining

## Lecture 8 - Recurrent Neural Networks (RNN)

Min Qian

Department of Biostatistics, Columbia University

October 24th, 2024

# Outline

- 1 Basic types of recurrent neural networks
- 2 Language modeling with RNN
- 3 RNN variations
- 4 Word embedding


## Examples of sequence data

## Speech recognition



“The quick brown fox jumped  
over the lazy dog.”

## Music generation

 $\emptyset$ 

## Sentiment classification

“There is nothing to like  
in this movie.”



## DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AGCCCCTGTGAGGAAGTAG

## Machine translation

Voulez-vous chanter avec  
moi?



Do you want to sing with  
me?

## Video activity recognition



## Running

## Name entity recognition

Yesterday, Harry Potter  
met Hermione Granger.




Yesterday, Harry Potter  
met Hermione Granger.  
Andrew Ng

# Name Entity Recognition


	Yesterday	Harry	Potter	met	Herminone	Granger.
Input X:	$X^{<1>}$	$X^{<2>}$	...	$X^{<t>}$	...	$X^{<T_x>}$
	0	1	1	0	1	1
Output Y:	$Y^{<1>}$	$Y^{<2>}$	...	$Y^{<t>}$	...	$Y^{<T_y>}$

Vocabulary of 10,000 words		Yesterday Harry Potter met Herminone Granger.					
One-hot encoding		$X^{<1>}$	$X^{<2>}$	$X^{<3>}$	$X^{<4>}$	$X^{<5>}$	$X^{<6>}$
	$\begin{pmatrix} a \\ aahed \\ \vdots \\ granger \\ \vdots \\ harry \\ \vdots \\ herminone \\ \vdots \\ met \\ \vdots \\ potter \\ \vdots \\ yesterday \\ \vdots \\ zyzzyva \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	

# Why Not Standard NN




Alternatively,




- Input/output can be of different lengths for different data points.
- Does not consider the sequential structure.

# Scheme of Basic RNN (Many-to-Many, Same Length)



From  $< t - 1 >$  to  $< t >$ :



Loss function:

$$l(Y, \hat{Y}) = \sum_{t=1}^{T_Y} l^{<t>} (Y^{<t>}, \hat{Y}^{<t>})$$


Parameters are learned using “backpropagation through time”

# Sentiment Classification: IMDb ratings (Many-to-One)


- Input: review

*This has to be one of the worst films of the 1990s. When my friends & I were watching this film (being the target audience it was aimed at) we just sat & watched the first half an hour with our jaws touching the floor at how bad it really was. The rest of the time, everyone else in the theater just started talking to each other, leaving or generally crying into their popcorn . . .*


- Output: sentiment of the review (positive/negative, rating 1-5).




# Summary of RNN types




One to one




One to many



Many to one



Many to many



Many to many

Andrew Ng



# Language Models

A **language model** is a probabilistic model of a natural language used to predict and generate plausible texts.

Speech recognition example:

- Sentence 1: She will bring flowers.
- Sentence 2: She will bring flours.

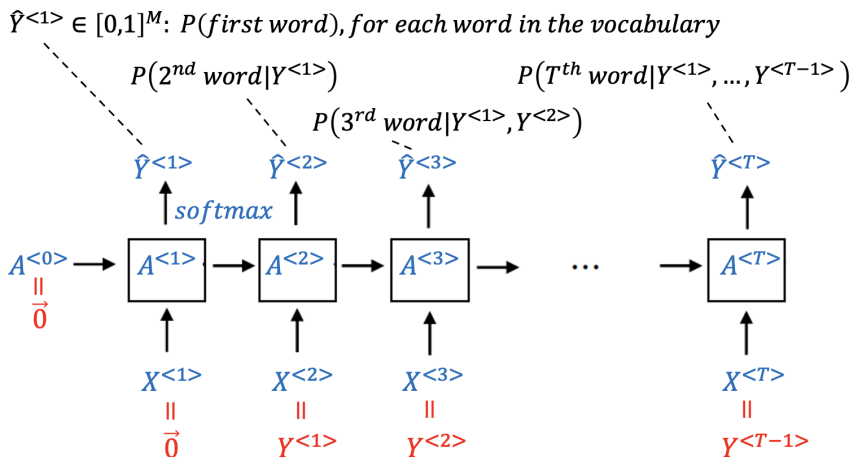
$P(sentence) = ?$

Training set: a large set (corpus) of English text.

She	will	bring	flowers	⟨EOS⟩
$Y^{<1>}$	$Y^{<2>}$	$Y^{<3>}$	$Y^{<4>}$	$Y^{<5>}$

- Can also add punctuation into the vocabulary dictionary.
- Add ⟨UNK⟩(unknown) into the vocabulary dictionary.


# RNN model



Loss function:  $L = \sum_{t=1}^T l(Y^{<t>}, \hat{Y}^{<t>})$ , where

$$l(Y^{<t>}, \hat{Y}^{<t>}) = - \sum_{m=1}^M Y_m^{<t>} \log(\hat{Y}_m^{<t>})$$

# Sampling a Sequence from a Trained RNN




- First, randomly sample a word from the distribution of  $\hat{Y}^{<1>}$ .
- For  $t = 2, 3, \dots$ , randomly sample a word using conditional distribution of  $\hat{Y}^{<t>}$  | previously sampled  $\hat{Y}^{<1>}, \dots, \hat{Y}^{<t-1>}$ .
- End the sentence if  $\langle \text{EOS} \rangle$  is sampled.
- If  $\langle \text{UNK} \rangle$  is sampled, can reject and resample.

# Long Term Dependency in RNN

The **dog**, after playing in the park for a long time ..., **sleeps** soundly.


The **dogs**, after playing in the park for a long time ..., **sleep** soundly.



- RNN is a very deep NN.
- $\hat{Y}^{<t>}$  mainly depends on near past information.
- Exploding gradients: gradient clipping. rescale gradient vectors if a gradient is above a threshold.
- Vanishing gradients?

# Gated Recurrent Unit (Simplified)


The **dog**, after playing in the park for a long time ..., **sleeps** soundly.



- **Update gate:**  $U^{<t>} = \text{sigmoid}(w_{UA}A^{<t-1>} + w_{UX}X^{<t>} + b_U)$
- **Current memory:**  $\tilde{A}^{<t>} = \tanh(w_{AA}A^{<t-1>} + w_{AX}X^{<t>} + b_A)$
- **Output:**  $A^{<t>} = U^{<t>} \odot \tilde{A}^{<t>} + (1 - U^{<t>}) \odot A^{<t-1>}$


# Gated Recurrent Unit (Simplified)

The **dog**, after playing in the park for a long time ..., **sleeps** soundly.



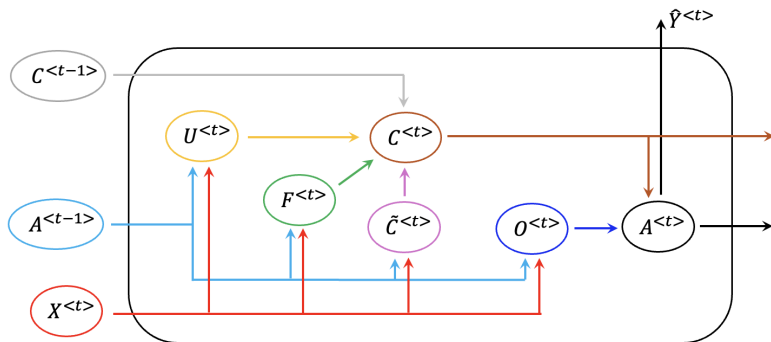
- **Update gate:**  $U^{<t>} = \text{sigmoid}(w_{UA}A^{<t-1>} + w_{UX}X^{<t>} + b_U)$
- **Current memory:**  $\tilde{A}^{<t>} = \tanh(w_{AA}A^{<t-1>} + w_{AX}X^{<t>} + b_A)$
- **Output:**  $A^{<t>} = U^{<t>} \odot \tilde{A}^{<t>} + (1 - U^{<t>}) \odot A^{<t-1>}$

# Full GRU (2014)



- **Update gate:**  $U^{<t>} = \text{sigmoid}(w_{UA}A^{<t-1>} + w_{UX}X^{<t>} + b_U)$
- **Reset gate:**  $\Gamma^{<t>} = \text{sigmoid}(w_{\Gamma A}A^{<t-1>} + w_{\Gamma X}X^{<t>} + b_{\Gamma})$
- **Current memory:**  
 $\tilde{A}^{<t>} = \tanh(w_{AA}(\Gamma^{<t>} \odot A^{<t-1>}) + w_{AX}X^{<t>} + b_A)$
- **Output:**  $A^{<t>} = U^{<t>} \odot \tilde{A}^{<t>} + (1 - U^{<t>}) \odot A^{<t-1>}$

# Long Short-Term Memory RNN (LSTM, 1997)



- **Update gate:**  $U^{<t>} = \text{sigmoid}(w_{UA}A^{<t-1>} + w_{UX}X^{<t>} + b_U)$
- **Forget gate:**  $F^{<t>} = \text{sigmoid}(w_{FA}A^{<t-1>} + w_{FX}X^{<t>} + b_F)$
- **Current memory:**  $\tilde{C}^{<t>} = \tanh(w_{CA}A^{<t-1>} + w_{CX}X^{<t>} + b_C)$
- **Memory cell:**  $C^{<t>} = U^{<t>} \odot \tilde{C}^{<t>} + F^{<t>} \odot C^{<t-1>}$
- **Output gate:**  $O^{<t>} = \text{sigmoid}(w_{OA}A^{<t-1>} + w_{OX}X^{<t>} + b_O)$
- **Output:**  $A^{<t>} = O^{<t>} \odot \tanh(C^{<t>})$




# GRU vs LSTM

- Both GRU and LSTM are designed to catch long-term dependency.
- LSTM was invented much earlier than GRU.
- GRU is a simpler model than LSTM. Computationally faster.
- GRU gains more momentum in recent years. It often just works as well as LSTM.

# Bi-directional RNN

Name entity recognition example:

- 1 I like **apple**, especially iphone.
- 2 I like apple. It is a healthy fruit.




- BRNN + LSTM is commonly used for NLP problem.
- Need entire sequence of data for prediction.

# Bi-directional RNN


Name entity recognition example:

- 1 I like **apple**, especially iphone.
- 2 I like apple. It is a healthy fruit.

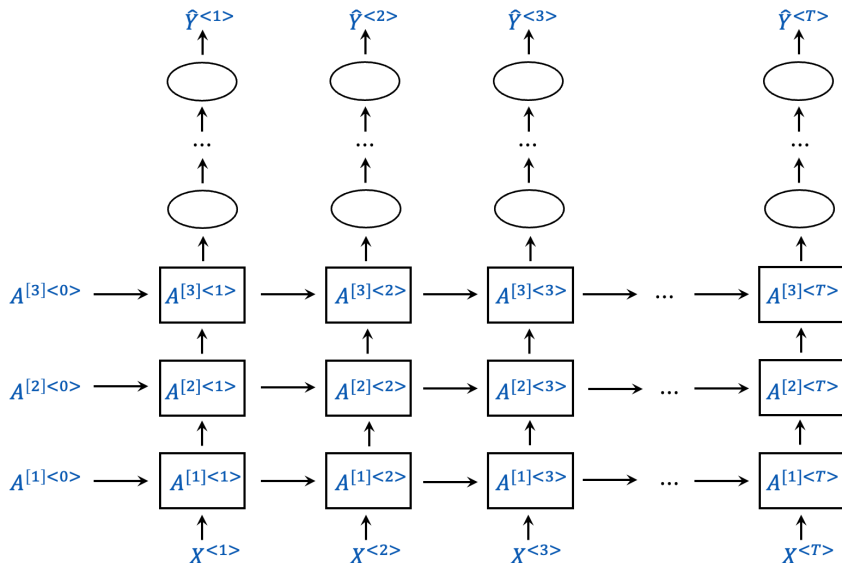


- BRNN + LSTM is commonly used for NLP problem.
- Need entire sequence of data for prediction.

# Deep RNN



# Most Commonly used Deep RNN



# Representing Words: One-hot Encoding


Each word is represented by a **one-hot-encoded** vector (i.e., dummy variable) with 9,999 zeros and a single 1 in some position.

Vocabulary of 10,000 words	man $x^{<1>}$	woman $x^{<2>}$	boy $x^{<3>}$	girl $x^{<4>}$	king $x^{<5>}$	queen $x^{<6>}$	apple $x^{<7>}$	orange $x^{<8>}$
$\begin{pmatrix} a \\ \vdots \\ apple \\ \vdots \\ boy \\ \vdots \\ girl \\ \vdots \\ king \\ \vdots \\ man \\ \vdots \\ orange \\ \vdots \\ queen \\ \vdots \\ woman \\ \vdots \\ zyzyzyva \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ \vdots \\ \vdots \\ 1 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ \vdots \\ \vdots \\ 1 \\ 0 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ 0 \end{pmatrix}$	

# Word Embedding: Featurized Representation

- The dimension of one-hot-encoding is very high.
- One-hot-encoding does not capture the similarity/correlation among words.

Representing each word (10,000-dim vector) by a  $K$ -dim (e.g., 300) vector.



Features	man	woman	boy	girl	king	queen	apple	orange
Masculinity	0.97	0.05	0.7	0.1	0.95	0.02	0.5	0.45
Age	0.75	0.7	0.25	0.23	0.6	0.58	0.49	0.51
Royalty	0.05	0.05	0.01	0.03	0.96	0.95	0.53	0.48
Edible	0.01	0.01	0.02	0.01	0.02	0.01	0.99	0.90
⋮								
⋮		⋮		⋮			⋮	
⋮		⋮		⋮			⋮	
Color		⋮		⋮			⋮	
Verb		⋮		⋮			⋮	
Size		⋮		⋮			⋮	

# Word Embedding: Embedding Matrix

Matrix  $\mathbf{E} \in \mathbb{R}^{K \times 10000}$ : maps a 10,000-dim vector to an  $K$ -dim vector.

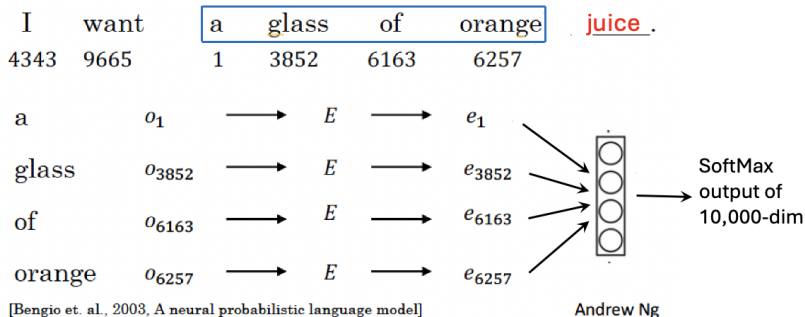
	10,000 columns							
	a	...	apple	...	...	orange	...	zyzzyva
	0.1125	...	0.5	...	...	0.45	...	0.9295
	0.7606	...	0.49	...	...	0.51	...	0.6156
	0.5513	...	0.53	...	...	0.48	...	0.7775
	0.6219	...	0.99	...	...	0.90	...	0.235
K rows	⋮			⋮				⋮
	⋮			⋮				⋮
	⋮			⋮				⋮
	⋮			⋮				⋮
	⋮			⋮				⋮
	⋮			⋮				⋮
	0.8586	...	0.15	...	...	0.13	...	0.9003

embedded vector  $e_{word} = \mathbf{E} \bullet \text{one-hot-vector } o_{word}$



# Learn Embedding Matrix using NN

Training set: a large set of **context (input)** / **target (output)** pairs



Various choices of context:

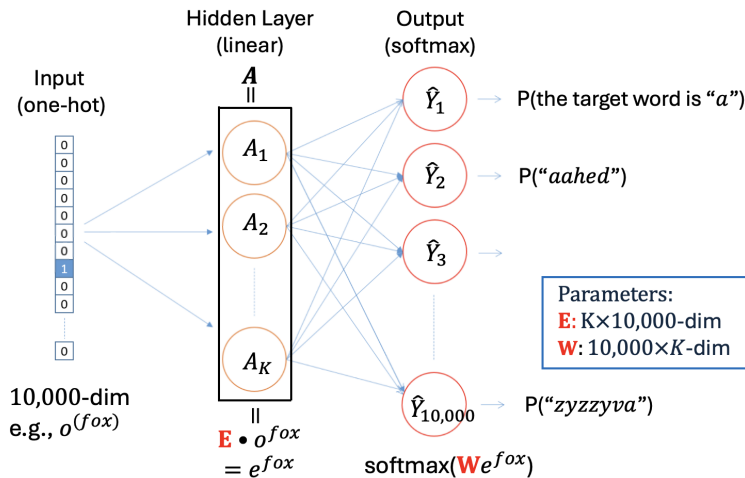
- Previous 4 words, previous 1 word, ...
- 4 words on left and 4 on right, e.g.,  
I want a glass of orange juice to pair with my meal.
- Nearby 1 word

# Word2Vec: Skip-Grams Training Data

- Idea: the words that appear in the same context (near each other) should have similar word vectors.
- Use an input word to predict surrounding words.
- Random pick a **context (input)** word and randomly pick a **target (output)** word within  $\pm x$  (e.g., 5) window of the context word.

Source Text	Training Samples						
<table><tr><td>The</td><td>quick</td><td>brown</td></tr></table> fox jumps over the lazy dog. ➡	The	quick	brown	(the, quick) (the, brown)			
The	quick	brown					
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td></tr></table> jumps over the lazy dog. ➡	The	quick	brown	fox	(quick, the) (quick, brown) (quick, fox)		
The	quick	brown	fox				
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td></tr></table> over the lazy dog. ➡	The	quick	brown	fox	jumps	(brown, the) (brown, quick) (brown, fox) (brown, jumps)	
The	quick	brown	fox	jumps			
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td><td>over</td></tr></table> the lazy dog. ➡	The	quick	brown	fox	jumps	over	(fox, quick) (fox, brown) (fox, jumps)
The	quick	brown	fox	jumps	over		

# The Skip-Gram Model



- The softmax step is very expensive to calculate
- Hierarchical classification is often used.
- Context words are sampled to balance common and less common words.

# Negative Sampling

e.g., The **quick** brown **fox** jumps over the lazy dog.

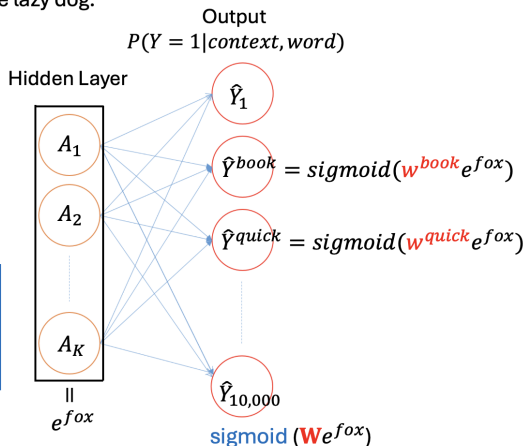
	context	word	target?
0	fox	quick	1
1	fox	book	0
2	fox	school	0
⋮	fox	phone	0
⋮	fox	pencil	0
$J$	fox	juice	0

Loss function:

$$\sum_{j=0}^J [-Y^{(j)} \log \hat{P}^{(j)} - (1 - Y^{(j)}) \log(1 - \hat{P}^{(j)})]$$

$$= -\log \hat{P}^{(0)} - \sum_{j=1}^J \log(1 - \hat{P}^{(j)})$$

Only involves rows in **W** for the  $J + 1$  words.



- For each positive example, randomly choose  $J$  negative words.  
 $J = 5 - 20$  for smaller datasets, and  $J = 2 - 5$  for larger datasets.

- sample negative target word  $w_i$  w.p.  $\frac{P(w_i)^{3/4}}{\sum_{j=1}^{10000} P(w_j)^{3/4}}$

# Global Vectors for Word Representation (GloVe)

- For each pair of words  $i$  and  $j$  in the vocabulary,  
 $X_{ij}$  = frequency that pair  $(i, j)$  is a pair of (context, target).
- Parameters:  
 $\mathbf{E} \in \mathbb{R}^{K \times 10,000}$ : each column  $e_i$  is a word vector.  
 $\mathbf{W} \in \mathbb{R}^{K \times 10,000}$ : each column  $w_j$  is another word vector.
- Estimate parameters  $\{(e_i, w_j, b_{e,i}, b_{w,j}) : i, j = 1, \dots, 10000\}$ :

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(e_i^\top w_j + b_{e,i} + b_{w,j} - \log X_{ij})^2,$$

- ▶  $f(0) = 0$ , by default  $0 \log 0 = 0$ .
  - ▶  $f(x)$  is non-decreasing
  - ▶  $f(x)$  is not too large for very large values of  $x$ .
- Final embedding vector for each word:  $(\hat{e}_i + \hat{w}_i)/2$ .

# Sentiment Classification: IMDb ratings



- Input: review

*This has to be one of the worst films of the 1990s. When my friends & I were watching this film (being the target audience it was aimed at) we just sat & watched the first half an hour with our jaws touching the floor at how bad it really was. The rest of the time, everyone else in the theater just started talking to each other, leaving or generally crying into their popcorn . . .*

- Response: sentiment of the review (positive/negative)

# Bag-of-words Model (ignore sequence of data)

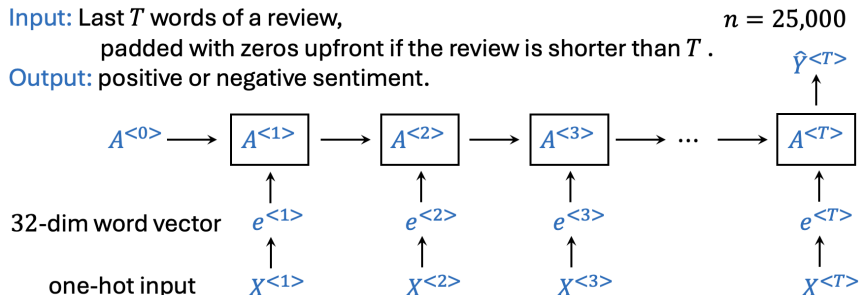
- Consider a vocabulary of  $M$  (say 10,000) words.
- For each review, the input is a binary vector of length 10,000 indicating whether a word is present or not.
- $n = 25,000$  for training and testing sets.



# IMDb ratings: “Entry Level” RNN models

**Input:** Last  $T$  words of a review,  
padded with zeros upfront if the review is shorter than  $T$ .

**Output:** positive or negative sentiment.



- Fit Simple RNN with  $K = 32$  hidden units, train embedding matrix. The test set accuracy is 76%.
- With LSTM RNN, the test set accuracy is increased to 87%.
- As of now, the leading RNN configurations report accuracy above 95% on the IMDb data.



# Summary

- RNN is a flexible tool for prediction with sequence input and/or output.
- GRU and LSTM are designed to catch long-term dependency.
- BRNN is designed to learn dependency of later information.
- Word embedding: map one-hot-vector to a dense vector.
  - ▶ learn using your training data
  - ▶ learn from a large text corpus.
  - ▶ use pre-trained embedding.

Reminder:

- Quiz for lecture 8 is due at 9pm on Monday Oct. 28th.
- Hw #2 is due at 9pm on Saturday Oct. 26th.