

P9120 - Statistical Learning and Data Mining

Lecture 11 - Reinforcement Learning III

Min Qian

Department of Biostatistics, Columbia University

November 14th, 2024

Outline

1 MDP planning problem

- Policy evaluation
- Policy improvement
- Policy iteration
- Value iteration

2 MDP learning problem

- TD(0) for policy evaluation
- Q-learning for estimating optimal policy

MDP Recap

$$S_0, A_0, R_0, \dots, S_t, A_t, R_t, \dots$$

- Markovian and stationary:

system dynamics: $p(s', r|s, a) = \Pr(S_{t+1} = s', R_t = r|S_t = s, A_t = a)$

state-transition: $p(s'|s, a) \triangleq \Pr(S_{t+1} = s'|S_t = s, A_t = a)$

exp. rewards for (s, a, s') : $r(s, a, s') \triangleq E(R_t|S_t = s, A_t = a, S_{t+1} = s')$

expected rewards for (s, a) : $r(s, a) \triangleq E(R_t|S_t = s, A_t = a)$

- **Return** (cumulative rewards) $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k}$.

- **Policy** Deterministic policy $\pi(s) : \mathcal{S} \rightarrow \mathcal{A}$;


Stochastic policy: $\pi(a|s) : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$;

- (state-)value function of π : $V_{\pi}(s) = E_{\pi}(G_t|S_t = s)$

- **Optimal policy** $\pi^*(s) = \arg \max_{\pi} V_{\pi}(s)$.

There always exists a deterministic optimal policy for an MDP.

Recycling Robot Example



- What's $V_{\pi}(s)$ of a given policy π ?

$$\pi(\text{high}) = \text{search}$$

$$\pi(\text{low}) = \text{wait}$$

Bellman Equation for Value Function $V_\pi(s)$

Note that $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k} = R_t + \gamma G_{t+1}$.

$$\begin{aligned} V_\pi(s) &\triangleq E_\pi[G_t | S_t = s] = E_\pi[R_t + \gamma G_{t+1} | S_t = s] \\ &= E_\pi[R_t + \gamma E_\pi(G_{t+1} | S_{t+1}, S_t = s) | S_t = s] \\ &= E_\pi[R_t + \gamma V_\pi(S_{t+1}) | S_t = s]. \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma V_\pi(s')] \\ &= \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a, s') + \gamma V_\pi(s')] \\ &= \sum_a \pi(a|s) \left[r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_\pi(s') \right] \end{aligned}$$

$V_\pi(s)$ is the unique solution to a system of linear equations!

Recycling Robot Example


$$V_{\pi}(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) \left[r(s, a, s') + \gamma V_{\pi}(s') \right]$$

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-

What is $V_{\pi}(s)$ for policy $\pi(\text{high}) = \text{search}, \pi(\text{low}) = \text{wait}$?

Gridworld Example

$$V_{\pi}(s) = \sum_a \pi(a|s) \left[r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_{\pi}(s') \right]$$



State space: 25 cells. **Action** space: $\leftarrow, \rightarrow, \uparrow, \downarrow$

Dynamics of MDP (reward dist. and transition prob.)

$S_{t+1} = S_t, R_t = -1$ if A_t would take the agent off the grid

$S_{t+1} = A', R_t = +10$ if $S_t = A$ regardless of A_t

$S_{t+1} = B', R_t = +5$ if $S_t = B$ regardless of A_t

$S_{t+1} = S_t + \text{one cell in the direction of } A_t, R_t = 0$, o.w.

Question: What is the value function of a random policy (i.e., $\pi(a|s) = 1/4$ for all (s, a) pairs) with discount rate $\gamma = 0.9$?

Planning: Policy Evaluation

Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$ arbitrarily, for $s \in \mathcal{S}$, and $V(\text{terminal})$ to 0

Loop:

$\Delta \leftarrow 0$



Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$



3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

Policy Improvement: Value Functions

- (state-)Value function

$$V_{\pi}(s) \triangleq E_{\pi}(G_0|S_0 = s) = E_{\pi}(G_t|S_t = s).$$

i.e., expected return if start in state s and then follow policy π .

- Action-value function (also called Q-function):

$$Q_{\pi}(s, a) \triangleq E_{\pi}(G_t|S_t = s, A_t = a).$$

expected return if start in s , take action a (once), then follow π .

- V_{π} can be computed from Q_{π} :

$$V_{\pi}(s) = E_{\pi}[Q_{\pi}(S_t, A_t)|S_t = s] = \sum_a \pi(a|s)Q_{\pi}(s, a).$$

- Bellman Equation for $Q_{\pi}(s, a)$

$$Q_{\pi}(s, a) = E_{\pi}[R_t + \gamma Q_{\pi}(S_{t+1}, A_{t+1})|S_t = s, A_t = a]$$

Policy Improvement Theorem

Let π and π' be a pair of policies such that, for all $s \in \mathcal{S}$,

$$Q_\pi(s, \pi'(s)) \geq V_\pi(s) \quad \text{if } \pi' \text{ is a deterministic policy,}$$

$$\text{or} \quad \sum_a \pi'(a|s) Q_\pi(s, a) \geq V_\pi(s) \quad \text{if } \pi' \text{ is a stochastic policy.}$$

Then $V_{\pi'}(s) \geq V_\pi(s)$ for all $s \in \mathcal{S}$.

Proof.

Note that $Q_\pi(s, \pi'(s)) = E[R_t + \gamma V_\pi(S_{t+1}) | S_t = s, A_t = \pi'(s)]$

$$\begin{aligned} V_\pi(s) &\leq E_{\pi'}[R_t + \gamma V_\pi(S_{t+1}) | S_t = s] \\ &\leq E_{\pi'}[R_t + \gamma E_{\pi'}[R_{t+1} + \gamma V_\pi(S_{t+2}) | S_{t+1}] | S_t = s] \\ &= E_{\pi'}[R_t + \gamma R_{t+1} + \gamma^2 V_\pi(S_{t+2}) | S_t = s] \\ &\leq E_{\pi'}[R_t + \gamma R_{t+1} + \gamma^2 E_{\pi'}[R_{t+2} + \gamma V_\pi(S_{t+3}) | S_{t+2}] | S_t = s] \\ &\leq E_{\pi'}[R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots | S_t = s] = V_{\pi'}(s) \end{aligned}$$



Planning: Policy Improvement

For a given policy π :

- 1 Compute Q_π from V_π :

$$\begin{aligned}Q_\pi(s, a) &= E_\pi[R_t + \gamma G_{t+1} | S_t = s, A_t = a] \\&= E_\pi[R_t + \gamma V_\pi(S_{t+1}) | S_t = s, A_t = a] \\&= \sum_{s', r} p(s', r | s, a) [r + \gamma V_\pi(s')] \\&= r(s, a) + \gamma \sum_{s'} p(s' | s, a) V_\pi(s')\end{aligned}$$

- 2 Improve π by the greedy policy $\pi'(s) = \arg \max_a Q_\pi(s, a)$. Since

$$Q_\pi(s, \pi'(s)) = \max_a Q_\pi(s, a) \geq Q_\pi(s, \pi(s)) = V_\pi(s),$$

we have $V_{\pi'}(s) \geq V_\pi(s)$ for all s .

Policy Iteration for Computing Optimal Policy π^*

- ① Initialization: arbitrary policy $\pi_0(s)$.
- ② For $t = 0, 1, 2, \dots$, repeat
 - ① Policy evaluation: Initialize $V_{\pi_t}^{(0)}(s)$ for $s \in \mathcal{S}$. For $j = 0, 1, 2, \dots$,

$$V_{\pi_t}^{(j+1)}(s) \leftarrow \sum_{s' \in \mathcal{S}} p(s'|s, \pi_t(s)) [r(s, \pi_t(s)) + \gamma V_{\pi_t}^{(j)}(s')] \text{ for each } s \in \mathcal{S}$$

until convergence. The converged value is denote as $V_{\pi_t}(s)$.

- ② If $\pi_t(s) = \arg \max_a \left[r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_{\pi_t}(s') \right]$ for $\forall s \in \mathcal{S}$, stop and output $\pi^* = \pi_t$ and $V^* = V_{\pi_t}$.

Otherwise, policy Improvement:

$$\pi_{t+1}(s) \leftarrow \arg \max_a \left[r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_{\pi_t}(s') \right].$$

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi^*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:



old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop and return $V \approx V^*$ and $\pi \approx \pi^*$; else go to 2


Bellman Equation for Value (Gridworld Example)



Gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

V^*



π^*

Optimal Value Functions

- Optimal Policy $\pi^*(s) = \arg \max_{\pi} V_{\pi}(s)$.
- Optimal value function

$$V^*(s) \triangleq \max_{\pi} V_{\pi}(s) = V_{\pi^*}(s).$$

- Optimal action-value function (or **optimal Q-function**)

$$\begin{aligned} Q^*(s, a) &\triangleq \max_{\pi} Q_{\pi}(s, a) \\ &= \max_{\pi} E[R_t + \gamma V_{\pi}(S_{t+1}) | S_t = s, A_t = a] \\ &= E[R_t + \gamma V_{\pi^*}(S_{t+1}) | S_t = s, A_t = a] \\ &= Q_{\pi^*}(s, a) \end{aligned}$$

Bellman Optimality Equations

$$\begin{aligned} V^*(s) &= E_{\pi^*}[R_t + \gamma V_{\pi^*}(S_{t+1}) | S_t = s] \\ &= Q_{\pi^*}(s, \pi^*(s)) = \max_a Q_{\pi^*}(s, a) \quad (1) \\ &= \max_a E[R_t + \gamma V^*(S_{t+1}) | S_t = s, A_t = a] \\ &= \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma V^*(s')] \end{aligned}$$

$$\begin{aligned} \text{and } Q^*(s, a) &= E[R_t + \gamma V^*(S_{t+1}) | S_t = s, A_t = a] \\ &= E[R_t + \gamma \max_a Q^*(S_{t+1}, a) | S_t = s, A_t = a] \\ &= \sum_{s'} p(s'|s, a) [r(s, a) + \gamma \max_a Q^*(s', a)] \end{aligned}$$

From (1), $\pi^*(s) = \arg \max_a Q^*(s, a)$.

Planning: Value Iteration

- 1 Find optimal value function using an iterative process based on bellman optimality equation.
- 2 Derive the optimal policy from the optimal value function

Value Iteration, for estimating $\pi \approx \pi^*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy, $\pi \approx \pi^*$, such that
$$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

The Learning Problem

- The environment dynamic is unknown, but data are available.
- How to learn $V_\pi(\cdot)$, i.e., the value function for a given policy π ?
 - Prediction problem
- How to learn the optimal policy π^* ?
 - Control problem

Bellman Equations:

$$V_\pi(S_t) = E_\pi [R_t + \gamma V_\pi(S_{t+1}) | S_t]$$
$$Q^*(S_t, A_t) = E [R_t + \gamma \max_a Q^*(S_{t+1}, a) | S_t, A_t].$$

Basic Optimization

Data (Y_1, \dots, Y_n) with $Y_i = \mu + \epsilon_i$ and $E\epsilon_i = 0$ for $i = 1, \dots, n$.

How to estimate μ ?

- Analytically: $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n Y_i = \arg \min_{\mu} \frac{1}{2} \sum_{i=1}^n (Y_i - \mu)^2$.
- Gradient descent: $\hat{\mu}^{new} \leftarrow \hat{\mu}^{old} + \alpha \sum_{i=1}^n (Y_i - \hat{\mu}^{old})$.
- Stochastic gradient descent:
update using one data point at a time $\hat{\mu}^{new} \leftarrow \hat{\mu}^{old} + \alpha(Y_i - \hat{\mu}^{old})$.

Bellman Equation: $V_{\pi}(S_t) = E_{\pi}[R_t + \gamma V_{\pi}(S_{t+1}) | S_t]$

Generate data based on policy π :

$$R_t + \gamma V_{\pi}(S_{t+1}) = V_{\pi}(S_t) + \epsilon_t, \text{ where } E(\epsilon_t | S_t) = 0.$$

TD(0) for Prediction (i.e., Policy Evaluation)

Given $S_t \rightarrow$ choose $A_t = \pi(S_t) \rightarrow$ observe (S_{t+1}, R_t) , update

$$\hat{V}_\pi^{new}(S_t) \leftarrow \hat{V}_\pi^{old}(S_t) + \alpha[R_t + \gamma V_\pi(S_{t+1}) - \hat{V}_\pi^{old}(S_t)].$$

Replace $R_t + \gamma V_\pi(S_{t+1})$ by $R_t + \gamma \hat{V}_\pi^{old}(S_{t+1})$.

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in S^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

Q-Learning (Watkins, 1989)

Bellman Equation for optimal Q-function

$$Q^*(S_t, A_t) = E[R_t + \gamma \max_a Q^*(S_{t+1}, a) | S_t, A_t]$$

implies that $R_t + \gamma \max_a Q^*(S_{t+1}, a) = Q^*(S_t, A_t) + \epsilon$, $E(\epsilon | S_t, A_t) = 0$.

- Off-policy: learn optimal policy while data collected from any policy.
- Useful if you cannot control the policy for data generation. (very useful in health application).
- Can reuse data generated from old policies.

Q-Learning (off-policy Control)

For any given (S_t, A_t, R_t, S_{t+1}) :

$$\hat{Q}^{new}(S_t, A_t) \leftarrow \hat{Q}^{old}(S_t, A_t) + \alpha [R_t + \gamma \max_a Q^*(S_{t+1}, a) - \hat{Q}^{old}(S_t, A_t)]$$

Replace $R_t + \gamma \max_a Q^*(S_{t+1}, a)$ by $R_t + \gamma \max_a \hat{Q}^{old}(S_{t+1}, a)$.

Q-learning (off-policy TD control) for estimating $\pi \approx \pi^*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Q-learning with Continuous State Space

- Bellman Equation $Q^*(S_t, A_t) = E[R_t + \gamma \max_a Q^*(S_{t+1}, a) | S_t, A_t]$.
- Approximate $Q^*(s, a)$ by $Q(s, a; \theta)$, and estimate θ by minimizing

$$\sum_t [R_t + \gamma \max_a Q(S_{t+1}, a; \theta) - Q(S_t, A_t; \theta)]^2.$$

- Stochastic gradient descent (viewing $Q(S_{t+1}, a; \theta)$ as a fixed target)

$$\hat{\theta}^{new} \leftarrow \theta + \alpha [R_t + \gamma \max_a Q(S_{t+1}, a; \theta) - Q(S_t, A_t; \theta)] \frac{dQ(S_t, A_t; \theta)}{d\theta} \Big|_{\theta=\hat{\theta}^{old}}.$$


	Q-learning
Tabular	Converges to q^* (Watkins and Dayan, 1992) (Tsitsiklis, 1994)
Linear	Can diverge (Wiering, 2004; Baird, 1995)
Non-Linear	Can Diverge

Deep Q-Network (DQN)

Data tuples $\{(S_t, A_t, R_t, S_{t+1}), t = 0, 1, 2, \dots\}$

Usually store 10,000 most recent (S_t, A_t, R_t, S_{t+1}) in replay memory \mathcal{D}

- 1 Initialize $Q(s, a)$ (say, $Q(s, a) = 0$ for all (s, a)).
- 2 Randomly sample a mini-batch of tuples (S_t, A_t, R_t, S_{t+1}) from \mathcal{D} and train the Q-function using neural network:



- 3 Set $Q = Q^{new}$

RL Mobile Health Applications

IntelliCare for Depression and Anxiety



The hub app:

- Recommend apps weekly
- Track app usage.
- Manage other apps and push notifications.

- What app(s) to recommend so as to maximize app usage?
 - ▶ app usage of current week
 - ▶ cumulative app usages over several weeks.
- When to send out push notification so the users will respond?

Reminders

- Final project proposal due at 9pm on December 2nd
 - ▶ Topic
 - ▶ The method you will investigate
 - ▶ Simulation, or real data analysis (what's the data set), or theoretical derivation of a novel method.
- Hw #3 is due at 9pm on November 16th
- In class group paper presentation (next week Nov 21st):
 - ▶ [U-Net: Convolutional Networks for Biomedical Image Segmentation](#)
Team member: Manye Dong, Jiatong Li, Yiming Li, Wenxin Tian, Yuntian Xu, Shihang Zeng
 - ▶ [Deep-learning-based real-time prediction of acute kidney injury outperforms human predictive performance.](#)
Team members: Ruoying Deng, Ruijie He, Ekaterina Hofrenning, Jessie Li, Authur Starodynov, Yueyi Xu